

[\(https://wordpress-](https://wordpress-courses2223.wolfware.ncsu.edu/csc-501-002-sprg-2023/)

CSC 501 (002) Spring 2023 Operating Systems Principles

[courses2223.wolfware.ncsu.edu/csc-501-002-sprg-2023/](https://wordpress-courses2223.wolfware.ncsu.edu/csc-501-002-sprg-2023/))

PA2

CSC 501: Spring 2023

PA2: Readers/Writer Locks with Priority Inheritance

Due: Friday, March 17, 2023, 11:59 PM

1. Introduction

In PA2, you are going to implement readers/writer locks as described in this handout. Additionally, you will implement a priority inheritance mechanism to prevent the priority inversion problem when using locks.

Please download and untar a fresh version of the XINU for QEMU source at ([csc501-lab2-qemu.tgz \(https://moodle-courses2223.wolfware.ncsu.edu/pluginfile.php/1237269/mod_assign/introattachment/0/csc501-lab2-qemu.tgz?forcedownload=1\)](https://moodle-courses2223.wolfware.ncsu.edu/pluginfile.php/1237269/mod_assign/introattachment/0/csc501-lab2-qemu.tgz?forcedownload=1)).

Readers/writer locks are used to synchronize access to a shared data structure. A lock can be acquired for read or write operations. A lock acquired for reading can be shared by other readers, but a lock acquired for writing must be exclusive.

You have been provided with the standard semaphore implementation for XINU. Make sure you read and understand the XINU semaphore system (`wait.c`, `signal.c`, `screate.c`, `sdelete.c`, etc.) and use that as a basis for your locks. You should NOT modify the standard semaphore implementation, since semaphores are used in the rest of the kernel, e.g., in device drivers. Your task is to extend the XINU semaphore to implement the readers/writer lock semantics. Although, the standard semaphores implemented in XINU are quite useful, there are some issues with the XINU semaphores which we will try to fix in this assignment.

XINU semaphores do not distinguish between read accesses, which can co-exist, and write accesses, which must be exclusive.

Another problem with XINU's semaphores occurs when a semaphore is deleted at a time when it has processes waiting in its queue. In such situation, *sdelete* awakens all the waiting processes by moving them from the semaphore queue to the ready list. As a result, a process that is waiting for some event to occur will be awakened, even though the event has not yet occurred. You need to fix this problem in this PA.

Yet another problem that occurs due to the interactions between process synchronization and process scheduling is priority inversion. Priority inversion occurs when a higher priority thread is blocked waiting on a lock (or a semaphore) held by a lower priority thread. This can lead to erroneous system behavior, especially in real time systems.

There are many solutions in the literature to solve the problem of priority inversion. In this lab, you will implement one such solution discussed in our lectures: *priority inheritance protocol for locks*. Other solutions, for example, can be found on [Wikipedia \(https://en.wikipedia.org/wiki/Priority_inversion\)](https://en.wikipedia.org/wiki/Priority_inversion).

2. Interfaces to Implement

Basic Locks

For this lab you must implement the entire readers/writer lock system. This includes code or functions to:

- > initialize locks (call a function `linit()` from the `sysinit()` function in `initialize.c`)
- > create and destroy a lock (`lcreate` and `ldelete`)
- > acquire a lock and release multiple locks (`lock` and `releaseall`)

Please create files called `linit.c`, `lcreate.c`, `ldelete.c`, `lock.c` and `releaseall.c` that contain these functions. Use a header file called `lock.h` for your definitions, including the constants `DELETED`, `READ` and `WRITE`. The functions have to be implemented as explained next:

- > Create a lock: `int lcreate (void)` – Creates a lock and returns a lock descriptor that can be used in further calls to refer to this lock. This call should return `SYSERR` if there are no available entries in the lock table. The number of locks allowed is `NLOCKS`, which you should define in `lock.h` to be 50.
- > Destroy a lock: `int ldelete (int lockdescriptor)` – Deletes the lock identified by the descriptor `lockdescriptor`. (see “Lock Deletion” below)
- > Acquisition of a lock for read/write: `int lock (int ldes1, int type, int priority)` – This call is explained below (“Wait on locks with Priority”).
- > Simultaneous release of multiple locks: `int releaseall (int numlocks, int ldes1, ..., int ldesN)`

(1) Lock Deletion

As mentioned before, there is a slight problem with XINU semaphores. The way XINU handles *sdelete* may have undesirable effects if a semaphore is deleted while a process or processes are waiting on it. Examining the code for `wait` and `sdelete`, you will notice that `sdelete` readies processes waiting on a semaphore being deleted. So they will return from `wait` with `OK`.

You must implement your lock system such that waiting on a lock will return a new constant `DELETED` instead of `OK` when returning due to a deleted lock. This will indicate to the user that the lock was deleted and not unlocked. As before, any calls to `lock()` after the lock is deleted should return `SYSERR`.

There is also another subtle but important point to note. Consider the following scenario. Let us say that there are three processes A, B, and C. Let A create a lock with descriptor=X. Let A and B use X to synchronize among themselves. Now, let us assume that A deletes the lock X. But B does not know about that. If, now, C tries to create a lock, there is a chance

that it gets the same lock descriptor as that of X (lock descriptors are limited and hence can be reused). When B waits on X the next time, it should get a `SYSERR`. It should not acquire the lock C has now newly created, even if this lock has the same id as that of the previous one. You have to find a way to implement this facility, in addition to the `DELETED` issue above.

(2) Locking Policy

In your implementation, no readers should be kept waiting unless (i) a writer has already obtained the lock, or (ii) there is a higher **or equal** lock priority writer already waiting for the lock. Hence, when a writer or the last reader releases a lock, the lock should be next given to a process having the highest lock priority for the lock. In the case of equal lock priorities among readers or writers, the lock will be first given to the reader or writer that has the longest waiting time (in milliseconds) on the lock. If a writer's lock priority is equal to the highest lock priority of the waiting reader, and the writer's waiting time is no more than 0.5 second longer (think of it as a grace period for writer), the writer should be given preference to acquire the lock over the waiting reader. In any case, if a reader is chosen to have a lock, then all the other waiting readers having priority **greater** than that of the highest-priority waiting writer for the same lock should also be admitted.

(3) Wait on Locks with Priority

This call allows a process to wait on a lock with priority. The call will have the form:

```
int lock (int ldes1, int type, int priority)
```

where priority is any integer priority value (including negative values, positive values and zero).

Thus when a process waits, it will be able to specify a wait priority. Rather than simply enqueueing the process at the end of the queue, the `lock()` call should now insert the process into the lock's wait list according to the wait priority. Please note that the wait priority is different from a process's scheduling priority specified in the `create(..)` system call. A larger value of the priority parameter means a higher priority.

Control is returned only when the process is able to acquire the lock. Otherwise, the calling process is blocked until the lock can be obtained.

Acquiring a lock has the following meaning:

1. The lock is free, i.e., no process is owning it. In this case the process that requested the lock gets the lock and sets the type of locking as `READ` or `WRITE`.
2. Lock is already acquired:
3. For `READ`:
If the requesting process has specified the lock type as `READ` and has sufficiently high priority (larger than the highest priority writer process waiting for the lock), it acquires the lock, else not.
4. For `WRITE`:
In this case, the requesting process does not get the lock as `WRITE` locks are exclusive.

(4) Releasing Locks

Simultaneous *release* allows a process to release one or more locks simultaneously. The system call has the form

```
int releaseall (int numlocks, int ldes1, ...)
```

and should be defined according to the locking policy given above. Also, each of the lock descriptors must correspond to a lock being held by the calling process.

If there is a lock in the arguments which is not held by calling process, this function needs to return SYSERR and should not release this lock. However, it will still release other locks which are held by the calling process.

(5) Using Variable Arguments

The call `releaseall (int numlocks,...)`, has a variable number of arguments. For instance, it could be:

```
releaseall(numlocks,ldes1, ldes2);
releaseall(numlocks,ldes1, ldes2, ldes3, ldes4);
```

where `numlocks = 2` in the first case and `numlocks = 4` in the second case.

The first call releases two locks `ldes1` and `ldes2`. The second releases four locks. You will not use the `va_list/va_arg` facilities to accommodate variable numbers of arguments, but will obtain the arguments directly from the stack. See `create.c` for hints on how to do this.

3. Priority Inheritance

Introduction

Note: The priority mentioned in this section is the process' scheduling priority and not the wait priority. The priority inheritance protocol solves the problem of priority inversion by increasing the priority of the low priority process holding the lock to the priority of the high priority process waiting on the lock.

Basically, the following invariant must be maintained for all processes p :

$\text{Prio}(p) = \max (\text{Prio}(p_i)), \quad \text{for all processes } p_i \text{ waiting on any of the locks held by process } p.$

Furthermore, you also have to ensure the transitivity of priority inheritance. This scenario can be illustrated with the help of an example. Suppose there are three processes A, B, and C with priorities 10, 20, and 30 respectively. Process A acquires a lock L1 and Process B acquires a lock L2. Process A then waits on the lock L2 and becomes ineligible for execution. If now process C waits on the lock L1, then the priorities of both the processes A and B should be raised to 30.

Priority Inheritance Implementation Hints

These hints give a possible implementation of the priority inheritance protocol. You are free to come up with your own implementation as long as the functionality is correct.

You can maintain the following information:

➤ Inside the process table entry:

1. a) The original priority (`pprio`) with which a process is created.
- b) `pinh` – The current inherited priority of the process. This value can be 0 when the process is running with its original priority.

- c) A bit mask or a linked list through which all the locks held by the process can be found.
- d) An integer value lockid indicating the lock descriptor in whose wait queue the process is blocked. It can take the value-1 when the process is not blocked inside any wait queue. Note that a process can be only inside a single wait queue at a time.

> Inside the lock descriptor table entry.

1. a) A priority field (lprio) indicating the maximum priority among all the processes waiting in the lock's wait queue.
- b) A linked list/bitmask of the process ids of the processes currently holding the lock.

The actions to be performed in the following situations are:

- > lock: Suppose that the process P1 requests a lock. If the lock is available, then nothing needs to be done. Otherwise, if the priority of the process (P2) holding the lock is no less than the priority of P1, then also nothing has to be done. Otherwise the priority of P2 (pinh) has to be ramped up to the priority of P1 (pinh if pinh != 0; pprio otherwise). Note that whenever we refer to the priority of a process, we refer to the pinh field if it is non-zero; pprio otherwise. After ramping up the priority of P1, we also have to take care of the transitivity of priority inheritance as discussed earlier.
- > Release: On releasing a lock, the priority of the process has to be reset to the maximum priority of all the processes in the wait queues of all the locks *still* held by the process. The bitmask/linked list field in the process table entry can be used for this purpose. Note that multiple locks can be released by a process.
- > chprio, kill, etc: These system calls can also have a side effect. If a process P1, in the wait queue of a lock L has its priority changed or is killed, we need to recalculate the maximum priority of all the processes in L's wait queue, and update the priority of the processes holding L too, if necessary.

This is a basic outline, but feel free to modify/improve upon this scheme. Questions should be posted to the message board.

You are provided with a code example in the file `test_main.c` (https://moodle-courses2223.wolfware.ncsu.edu/pluginfile.php/1237269/mod_assign/introattachment/0/test_main.c?forcedownload=1).

4. Additional Tasks

Write your answers to the following questions in a file named Lab2Answers.txt(in simple text).

Please place this file in the TMP/ directory and turn it in, along with the above programming assignment.

1. The priority inversion problem could cause that a higher-priority process needs to wait for the completion of a lower-priority process.
 1. Priority inheritance is one approach to address this issue. Please give another approach and briefly describe its idea.
 2. Design a test case to show that your approach can address the priority inversion problem. You should compare the results that derived from the original XINU implementation (use semaphore) and that derived from your implementation (use readers/writer locks). Put your results in Lab2Answers.txt and your test case program (name it task1.c) in both ./sys and ./TMP
2. Synchronization is a quite complex issue in operating systems. To show your understanding about process synchronization, you are asked to point out a potential problem in this pseudo-code: `task2_sync_issue.c` (https://moodle-courses2223.wolfware.ncsu.edu/pluginfile.php/1237269/mod_assign/introattachment/0/task2_sync_issue.c?forcedownload=1). Multiple reader processes can execute this code concurrently by calling reader(). You are asked to describe an issue it may cause and why it can happen. Put your answer in Lab2Answers.txt mentioned above.

Turn-in Instructions

Electronic turn-in instructions

1. go to the csc501-lab2-qemu/compile directory and do “make clean”.
2. go to the directory of which your csc501-lab2-qemu directory is a subdirectory (NOTE: please do not rename csc501-lab2-qemu, or any of its subdirectories.) e.g., if /home/csc501/csc501-lab2-qemu is your directory structure, goto /homes/csc501
3. create a subdirectory TMP (under the directory csc501-lab2-qemu) and copy all the files you have modified/written, both .c files and .h files into the directory.
4. compress the csc501-lab2-qemu directory into a tgz file and submit your file on Moodle. Please only upload one tgz file.

```
tar czf csc501-lab2.tgz csc501-lab2-qemu
```

Also, please remember that you are allowed and encouraged to share test cases on Piazza.

Make sure that your code compiles, debugging output are turned off, turn-in instructions are followed precisely. Verify your files submitted as no exceptions will be made for accidentally submitted wrong files. Submitting file with wrong format or wrong name will be treated as “source code cannot be compiled”.

Grading Policy

- (30%) Source code can be compiled and the generated image is bootable
- (60%) All test cases worth 60 points (it might include 8-12 test cases)
- (10%) Each additional question earns 5 points

Note

- POST your questions on [Piazza](https://piazza.com/class/lcl245vvyycu3ej) (<https://piazza.com/class/lcl245vvyycu3ej>). Please do not send email if not necessary
- Welcome to share your test cases on the message board, but not any programming details
- Test cases available only for reference but not for correctness:
 - [test_main.c](https://moodle-courses2223.wolfware.ncsu.edu/pluginfile.php/1237269/mod_assign/introattachment/0/test_main.c?forcedownload=1) (https://moodle-courses2223.wolfware.ncsu.edu/pluginfile.php/1237269/mod_assign/introattachment/0/test_main.c?forcedownload=1).

Recent Comments

Archives

Categories

> No categories

© 2023 NC State University. All rights reserved. [Accessibility \(http://accessibility.ncsu.edu/\)](http://accessibility.ncsu.edu/) [Privacy \(http://www.ncsu.edu/privacy/\)](http://www.ncsu.edu/privacy/) [University Policies \(http://policies.ncsu.edu/\)](http://policies.ncsu.edu/) [Login \(https://wordpress-courses2223.wolfware.ncsu.edu/csc-501-002-sprg-2023/wp-login.php\)](https://wordpress-courses2223.wolfware.ncsu.edu/csc-501-002-sprg-2023/wp-login.php)